

# NYC Taxi Fare Predictor

September 25, 2018

## 1 CSE 519 -- Data Science (Fall 2018)

## 2 Homework 2: Exploratory Data Analysis in iPython

### 2.1 New York City Taxi Fare Prediction

### 2.2 Md Majid Jahangir - SBU ID 112077145

#### 2.2.1 Setting up the libraries and packages and global variables

```
In [ ]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from geopy import distance
from sklearn.metrics.pairwise import paired_distances
from scipy.stats.stats import pearsonr
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from math import sqrt
plt.style.use('seaborn-whitegrid')

trainingDF = None      #training data
testDF = None          #test data
validateData = None    #validate data
```

#### 2.2.2 Methods to load training, test and validate data

```
In [ ]: #method to load train data in dataframe
def loadTrainingData():
    global trainingDF
    trainingDF = pd.read_csv('train.csv',nrows=(5000000+50000))

#method to load data to get error metrics, 50000 rows of the 5000000 train dataset
def loadValidateData():
    global validateData
    global trainingDF
```

```

        validateData = trainingDF[5000001:]
        trainingDF = trainingDF[:5000000]

        #method to load test data
        def loadTestData():
            global testDF
            testDF = pd.read_csv('test.csv')

In [ ]: loadTrainingData()
        loadTestData()
        loadValidateData()

```

## 2.2.3 Describing on test and training data to look at some statistics

```

In [4]: testDF.describe()
        trainingDF.describe()

```

```

Out[4]:
      fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude \
count  5.000000e+06      5.000000e+06      5.000000e+06      4.999964e+06
mean    1.134080e+01     -7.250678e+01      3.991974e+01     -7.250652e+01
std     9.820175e+00      1.280970e+01      8.963509e+00      1.284777e+01
min    -1.000000e+02     -3.426609e+03     -3.488080e+03     -3.412653e+03
25%     6.000000e+00     -7.399206e+01      4.073491e+01     -7.399139e+01
50%     8.500000e+00     -7.398181e+01      4.075263e+01     -7.398016e+01
75%     1.250000e+01     -7.396711e+01      4.076712e+01     -7.396367e+01
max      1.273310e+03      3.439426e+03      3.310364e+03      3.457622e+03

      dropoff_latitude  passenger_count
count  4.999964e+06      5.000000e+06
mean    3.991725e+01      1.684695e+00
std     9.486767e+00      1.331854e+00
min    -3.488080e+03      0.000000e+00
25%     4.073404e+01      1.000000e+00
50%     4.075315e+01      1.000000e+00
75%     4.076811e+01      2.000000e+00
max      3.345917e+03      2.080000e+02

```

```

In [5]: trainingDF.head(2)      #checking the format of data

```

```

Out[5]:
      key  fare_amount  pickup_datetime \
0  2009-06-15 17:26:21.0000001      4.5  2009-06-15 17:26:21 UTC
1  2010-01-05 16:52:16.0000002     16.9  2010-01-05 16:52:16 UTC

      pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude \
0      -73.844311      40.721319      -73.841610      40.712278
1      -74.016048      40.711303      -73.979268      40.782004

      passenger_count
0                  1
1                  1

```

## 2.3 Cleaning of Data

```
In [6]: def drop_null_values(df):  
        print(df.isnull().sum())      #finding any NaN or null values in the data  
        df = df.dropna(how = 'any', axis = 'rows')      #dropping off rows with NaN values
```

2.3.1 Also, from train data describe, we saw the max passenger count as great as 208. As per NYC Taxi and Limousine Commission rule, max 6 people can board a taxi at a time. Hence dropping any rows with count > 6 and < 1

```
In [7]: def clean_passenger_column(df):  
        df = df.drop(df[df['passenger_count']<1].index, axis = 0)  
        df = df.drop(df[df['passenger_count']>6].index, axis = 0)
```

2.3.2 Adding two features 'absolute difference of latitude' and 'absolute difference of longitude'

```
In [8]: #finding Manhattan vector and adding as a feature in the dataframe to find possible outliers  
def add_Manhattan_vector_feature(trainingDF):  
    trainingDF['abs_diff_long'] = (trainingDF.dropoff_longitude - trainingDF.pickup_longitude).abs()  
    trainingDF['abs_diff_lat'] = (trainingDF.dropoff_latitude - trainingDF.pickup_latitude).abs()
```

2.3.3 dropping off rows which has absolute difference of latitude more than 4.2 as New York state width is approx. 285 miles dropping off rows which has absolute difference of longitude more than as New York state length is approx. 330 miles as one latitude and longitude difference equals ~69 and ~55 miles respectively

```
In [9]: def remove_outliers_from_manhattan_distance(df):  
        df = df.drop(df[df['abs_diff_long']>=5.0].index, axis = 0)  
        df = df.drop(df[df['abs_diff_lat']>=4.2].index, axis = 0)
```

2.3.4 Bounding our train data by the max. and min. bounds of pick up and drop off latitudes and longitudes of test data as that's the range with which we will check against.

```
In [10]: def bound_train_data_basedon_test_data():  
        test_data_lat_range= min(testDF.pickup_latitude.min(), testDF.dropoff_latitude.min())  
        test_data_long_range = min(testDF.pickup_longitude.min(), testDF.dropoff_longitude.min())  
        test_data_range = (test_data_long_range, test_data_lat_range)  
        return test_data_range
```

```
In [11]: #bounding on train data using the above bounds  
def bound_data_based_on_test_data():  
    global trainingDF  
    global validateData  
    test_data_range = bound_train_data_basedon_test_data()  
    trainingDF = trainingDF[(trainingDF['pickup_latitude']>=test_data_range[1][0]) &  
                            (trainingDF['dropoff_latitude']>=test_data_range[1][0]) & (trainingDF['pickup_longitude']<=test_data_range[0][1]) &  
                            (trainingDF['dropoff_longitude']<=test_data_range[0][1]) & (trainingDF['pickup_latitude']<=test_data_range[1][0]) &  
                            (trainingDF['dropoff_latitude']<=test_data_range[1][0]) & (trainingDF['pickup_longitude']>=test_data_range[0][1]) &  
                            (trainingDF['dropoff_longitude']>=test_data_range[0][1])]
```

```

        (trainingDF['pickup_longitude']>=test_data_range[0][0]))
validateData = validateData[(validateData['pickup_latitude']>=test_data_range[1][0]
    (validateData['dropoff_latitude']>=test_data_range[1][0]) & (validateData['pickup_longitude']<=test_data_range[0][1]) &
    (validateData['dropoff_longitude']<=test_data_range[0][1]) & (validateData['pickup_longitude']>=test_data_range[0][0]))]

```

```

In [12]: #let's us describe the data again
trainingDF.describe()

```

```

Out[12]:      fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude  \
count    5.000000e+06      5.000000e+06      5.000000e+06      4.999964e+06
mean      1.134080e+01     -7.250678e+01      3.991974e+01     -7.250652e+01
std        9.820175e+00      1.280970e+01      8.963509e+00      1.284777e+01
min       -1.000000e+02     -3.426609e+03     -3.488080e+03     -3.412653e+03
25%        6.000000e+00     -7.399206e+01      4.073491e+01     -7.399139e+01
50%        8.500000e+00     -7.398181e+01      4.075263e+01     -7.398016e+01
75%        1.250000e+01     -7.396711e+01      4.076712e+01     -7.396367e+01
max         1.273310e+03      3.439426e+03      3.310364e+03      3.457622e+03

      dropoff_latitude  passenger_count
count    4.999964e+06      5.000000e+06
mean      3.991725e+01      1.684695e+00
std        9.486767e+00      1.331854e+00
min       -3.488080e+03      0.000000e+00
25%        4.073404e+01      1.000000e+00
50%        4.075315e+01      1.000000e+00
75%        4.076811e+01      2.000000e+00
max         3.345917e+03      2.080000e+02

```

### 2.3.5 Adding euclidean distance feature to the dataset

```

In [13]: def add_euclidean_distance_feature(df):
    geoloc_pickup_pair = list(zip(df.pickup_latitude, df.pickup_longitude))
    geoloc_dropoff_pair = list(zip(df.dropoff_latitude, df.dropoff_longitude))
    df['euc_distance'] = paired_distances(geoloc_pickup_pair, geoloc_dropoff_pair)

```

2.3.6 From data describe we saw there were negative fares and also, we know NY lies above equator and west of Prime meridian because of which the valid latitude and longitude range becomes (0,90) and (-180,0) respectively. We drop off such rows with invalid range of data.

```

In [14]: def drop_all_outliers_of_params(df):
    df = df.drop(((df[df.fare_amount<0]) | (df[df.pickup_latitude<0]) | (df[df['pickup_

```

### 2.3.7 Adding haversine distance feature to train data

```

In [15]: #method to compute haversine distance in miles between 2 coordinates - taken from sta
def haversine_distance(lat1, long1, lat2, long2,df):

```

```

data = [df]
for i in data:
    p = 0.017453292519943295
    a = 0.5 - np.cos((i[lat2] - i[lat1]) * p)/2 + np.cos(i[lat1] * p) * np.cos(i[
    i['haversine_distance'] = 0.6213712 * 12742 * np.arcsin(np.sqrt(a))

```

### 2.3.8 Parsing 'pickup\_datetime' field as a date field and adding time features

```

In [16]: def parse_dates_and_add_time_features(df,field,format1):
df[field] = pd.to_datetime(df[field],format=format1)
train_data = [df]
for row in train_data:
    row['year'] = row[field].dt.year
    row['month'] = row[field].dt.month
    row['date'] = row[field].dt.day
    row['day_of_week'] = row[field].dt.dayofweek
    row['hour'] = row[field].dt.hour
    row['quarter'] = row[field].dt.quarter
    row['sec'] = row[field].dt.second

```

### 2.3.9 Removing rows which has fare lesser than 2.5 USD because as per NYC Taxi and Limousine Commission that is the minimum base fare available.

```

In [17]: trainingDF=trainingDF[trainingDF.fare_amount>=2.5]      #todo

```

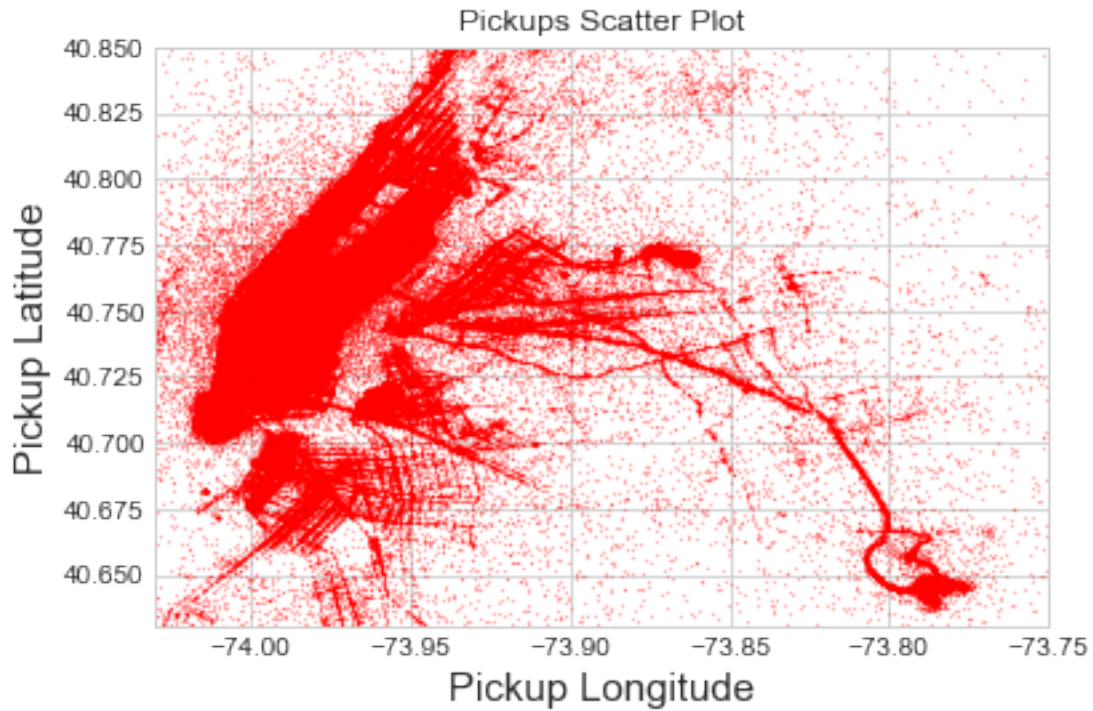
### 2.3.10 Creating a pickup heatmap scatter plot to visualise the hotspots for pickup

```

In [54]: plt.figure(figsize=(50,30))
trainingDF.plot(kind='scatter', x='pickup_longitude', y='pickup_latitude',
                color='red',
                s=.03, alpha=.7)
plt.title("Pickups Scatter Plot")
long_border = (-74.03, -73.75)
lat_border = (40.63, 40.85)
plt.ylim(lat_border)
plt.xlim(long_border)
plt.xlabel('Pickup Longitude', size=15)
plt.ylabel('Pickup Latitude',size=15)
plt.savefig('pickup_scatter.png')

```

<Figure size 3600x2160 with 0 Axes>



### 2.3.11 Generating features classifying data entries as pickup and drop off regions for 3 airports in NY

```
In [41]: airports={'JFK':{'min_long':-73.8352,
    'min_lat':40.6195,
    'max_long':-73.7401,
    'max_lat':40.6659},

    'EWR':{'min_long':-74.1925,
    'min_lat':40.6700,
    'max_long':-74.1531,
    'max_lat':40.7081
    },
    'LaG':{'min_long':-73.8895,
    'min_lat':40.7664,
    'max_long':-73.8550,
    'max_lat':40.7931
    }

}

def isAirport(lat,long,airport_name='JFK'):
    latMinTrue = (lat>=airports[airport_name]['min_lat'] )
```

```

latMaxTrue = (lat<=airports[airport_name]['max_lat'])
longMinTrue = (long>=airports[airport_name]['min_long'])
longMaxTrue = (long<=airports[airport_name]['max_long'])
#result = np.where( ((latitude>=nyc_airports[airport_name]['min_lat']) & (latitude
if latMinTrue and latMaxTrue and longMinTrue and longMaxTrue:
    return 1
else:
    return 0
return 0

def categorise_airport_feature(df):
    data = [df]
    for index,row in df.iterrows():
        row['is_pickup_JFK']=isAirport(row['pickup_latitude'],row['pickup_longitude'])
        row['is_dropoff_JFK']=isAirport(row['dropoff_latitude'],row['dropoff_longitude'])
        row['is_pickup_EWR']=isAirport(row['pickup_latitude'],row['pickup_longitude'])
        row['is_dropoff_EWR']=isAirport(row['dropoff_latitude'],row['dropoff_longitude'])
        row['is_dropoff_lag']=isAirport(row['dropoff_latitude'],row['dropoff_longitude'])
        row['is_pickup_lag']=isAirport(row['pickup_latitude'],row['pickup_longitude'])

```

### 2.3.12 Cleaning all the data based on above stated factors

```

In [20]: drop_null_values(trainingDF)
         drop_null_values(validateData)
         clean_passenger_column(trainingDF)
         clean_passenger_column(validateData)
         drop_all_outliers_of_params(validateData)
         bound_data_based_on_test_data()

```

```

key          0
fare_amount  0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude  36
dropoff_latitude  36
passenger_count  0
dtype: int64

key          0
fare_amount  0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude  0
dropoff_latitude  0
passenger_count  0
dtype: int64

```

### 2.3.13 Adding features described above

```
In [ ]: add_Manhattan_vector_feature(trainingDF)
        add_Manhattan_vector_feature(validateData)
        add_Manhattan_vector_feature(testDF)
        add_euclidean_distance_feature(trainingDF)
        add_euclidean_distance_feature(testDF)
        add_euclidean_distance_feature(validateData)
        drop_all_outliers_of_params(trainingDF)
        haversine_distance('pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude', trainingDF)
        haversine_distance('pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude', validateData)
        haversine_distance('pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude', testDF)
        parse_dates_and_add_time_features(trainingDF, 'pickup_datetime', "%Y-%m-%d %H:%M:%S UTC")
        parse_dates_and_add_time_features(testDF, 'pickup_datetime', "%Y-%m-%d %H:%M:%S UTC")
        parse_dates_and_add_time_features(validateData, 'pickup_datetime', "%Y-%m-%d %H:%M:%S UTC")

        #Adding feature column
        trainingDF['is_pickup_JFK'] = 0
        trainingDF['is_dropoff_JFK'] = 0
        trainingDF['is_pickup_EWR'] = 0
        trainingDF['is_dropoff_EWR'] = 0
        trainingDF['is_dropoff_lag'] = 0
        trainingDF['is_pickup_lag'] = 0
        testDF['is_pickup_JFK'] = 0
        testDF['is_dropoff_JFK'] = 0
        testDF['is_pickup_EWR'] = 0
        testDF['is_dropoff_EWR'] = 0
        testDF['is_dropoff_lag'] = 0
        testDF['is_pickup_lag'] = 0
        validateData['is_pickup_JFK'] = 0
        validateData['is_dropoff_JFK'] = 0
        validateData['is_pickup_EWR'] = 0
        validateData['is_dropoff_EWR'] = 0
        validateData['is_dropoff_lag'] = 0
        validateData['is_pickup_lag'] = 0
        categorise_airport_feature(testDF)
        categorise_airport_feature(trainingDF)
        categorise_airport_feature(validateData)
```

### 2.3.14 Using external dataset to add a feature 'is\_holiday'

```
In [ ]: #Holiday for USA dataset from 2010-2020. Source - https://data.world/jennifer-v/us-holiday
        holidayList = pd.read_csv('holidaydatelist.csv')
        holidayList.describe()
        holidayList.head(2)
        holidayList = holidayList[['day_date']]
        parse_dates_and_add_time_features(holidayList, 'day_date', "%Y-%m-%d")
        holidayList = holidayList.dropna(how = 'any', axis = 'rows')
```

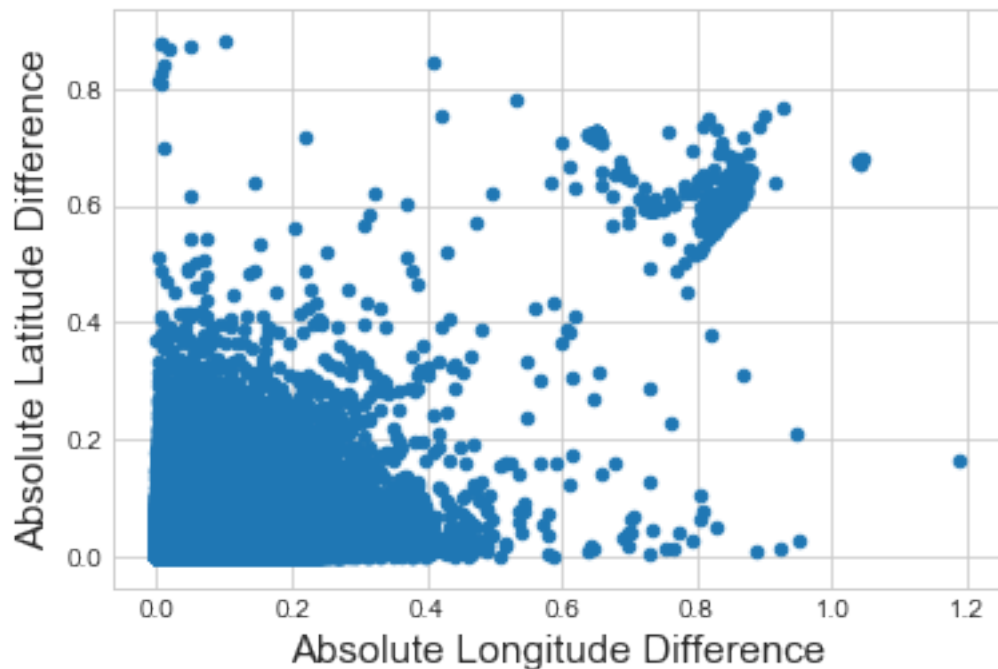


```
def add_is_holiday_feature(df1,df2):
    df1['is_holiday'] = 0
    date = list(df2.date)
    year = list(df2.year)
    month = list(df2.month)
    for index,row in df1.iterrows():
        for i in range(len(date)):
            if row.date==date[i] and row.month==month[i]:
                row['is_holiday'] = 1
add_is_holiday_feature(trainingDF,holidayList)
add_is_holiday_feature(testDF,holidayList)
add_is_holiday_feature(validateData,holidayList)
```

```
In [55]: #plotting the scatter plot to see if there are any outliers
plt.figure(figsize=(50,30))
plot = trainingDF.iloc[:].plot.scatter('abs_diff_long', 'abs_diff_lat')
plt.xlabel('Absolute Longitude Difference', size=15)
plt.ylabel('Absolute Latitude Difference',size=15)
```

```
Out[55]: Text(0,0.5,'Absolute Latitude Difference')
```

<Figure size 3600x2160 with 0 Axes>



```
In [23]: remove_outliers_from_manhattan_distance(trainingDF)
         remove_outliers_from_manhattan_distance(validateData)
```

**2.3.15** Since dataset is huge, let's try our first model by directly removing any distance value = 0, then later in further models we will try to impute those 0 values if it will be possible



```
In [24]: trainingDF = trainingDF.drop(trainingDF[trainingDF['haversine_distance']<=0].index, axis=1)
In [25]: #method to compute Pearson coefficient
def find_Pearson_coeeficent(df,x,y):
    corr, p_value = pearsonr(df[x], df[y])
    return corr
In [26]: find_Pearson_coeeficent(trainingDF, 'hour','fare_amount') #Solution to 2.c of
Out[26]: -0.017046104190507432
```

**2.3.16** Plotting graphs to show relation between the above 2 variables

```
In [53]: plt.figure(figsize=(25,7))
#plot = trainingDF.iloc[:].plot.scatter('haversine_distance', 'fare_amount')
#plt.rc('font', family='serif', size=1)
aggregate=trainingDF.groupby(['hour'])['fare_amount'].agg('mean')
plt.subplot(1, 2, 1)
gr = aggregate.plot(kind='bar')
gr.set_ylabel("Average Fare in USD", size = 20)
gr.set_xlabel("Time of the day in hour",size=20)

plt.subplot(1, 2, 2)

#plt.rc('font', family='serif', size=40)
plt.scatter(x=trainingDF['hour'], y=trainingDF['fare_amount'], s=1.5)
plt.xlabel('Time of the day - In hour', size=20)
plt.ylabel('Fare in USD',size=20)
plt.savefig('fare vs time of day.png')
```





```
In [29]: find_Pearson_coeeficent(trainingDF, 'fare_amount','haversine_distance')
Out[29]: 0.8333702568239763
In [60]: find_Pearson_coeeficent(trainingDF, 'hour','euc_distance')
Out[60]: -0.028565521465992692
```

### 2.3.17 Plotting the chart to visualise relation between hour and euclidean distance

```
In [63]: plt.figure(figsize=(25,7))
         aggregate=trainingDF.groupby(['hour'])['euc_distance'].agg('mean')
         plt.subplot(1, 2, 1)
         gr = aggregate.plot(kind='bar')
         gr.set_ylabel("Euclidean Distance", size = 20)
         gr.set_xlabel("Time of the day in hour",size=20)


         plt.subplot(1, 2, 2)

         #plt.rc('font', family='serif', size=40)
         plt.scatter(x=trainingDF['hour'], y=trainingDF['euc_distance'], s=1.5)
         plt.xlabel('Time of the day - In hour', size=20)
         plt.ylabel('Euclidean Distance',size=20)
         plt.savefig('distance vs time.png')
```



```
In [46]: #Plotting relation between distance and fare amount
         plt.figure(figsize=(50,30))
         plt.scatter(x=trainingDF['haversine_distance'], y=trainingDF['fare_amount'], s=1.5)
         plt.xlabel('Distance', size = 40)
         plt.ylabel('Fare', size = 40)
```

```
Out[46]: Text(0,0.5,'Fare')
```



```
In [27]: find_Pearson_coeeficent(trainingDF, 'hour','euc_distance') #Solution to 2.b o
```

```
Out[27]: -0.028565521465992692
```


```
In [28]: find_Pearson_coeeficent(trainingDF, 'euc_distance','fare_amount') #Solution to 2.a o
```

```
Out[28]: 0.8415478510506471
```

**2.3.18** Seems to be a linear relation with fares, with fares increasing with distance, also there are some outliers in fares which we should take care of.

```
In [57]: #plotting the relation between euclidean distance and fare_amount feature
```

```
plt.scatter(trainingDF['euc_distance'][:1000], trainingDF['fare_amount'][:1000])
plt.xlabel('Euclidean distance', size=15)
plt.ylabel('Fare $USD', size=15)
plt.show()
plt.savefig('euclidean_Distance_vs_fare_amount.png')
```



<Figure size 432x288 with 0 Axes>

## 2.4 Modelling our data

### 2.4.1 Basic Linear regression Model

```
In [58]: #creating input features to feed into model
def create_input_features(df):
    return np.column_stack((df.is_pickup_JFK, df.is_dropoff_JFK, df.is_dropoff_EWR, df

trainingData_X = create_input_features(trainingDF)
trainingData_Y = np.array(trainingDF.fare_amount)

In [59]: #taken from Kaggle starter notebook
(weight, _, _, _) = np.linalg.lstsq(trainingData_X, trainingData_Y, rcond = None)
print(weight)

[ 0.00000000e+00  5.54334356e-13 -2.04725126e-12  1.10422782e-12
  1.23989707e-12 -6.76347867e-13  3.40672560e+00 -2.00565423e+01
 -1.50680011e+00 -8.64900760e+00  3.85974645e-02 -1.17641164e+02
 -2.53573743e+02  7.19151321e+00 -5.12348803e-02  1.23855793e-02
  5.24400463e-01  2.59764048e+02]
```

## 2.4.2 Running the model

```
In [64]: test_X = create_input_features(testDF)
         # Predict fare_amount on the test set using our model (w) trained on the training set
         test_y_predictions = np.matmul(test_X, weight).round(decimals = 2)

In [ ]: # Write the predictions to a CSV file which we can submit to the competition.

def write_to_submission_csv(prediction):
    global testDF
    submission = pd.DataFrame(
        {'key': testDF.key, 'fare_amount': prediction},
        columns = ['key', 'fare_amount'])
    submission.to_csv('submission.csv', index = False)

In [ ]: write_to_submission_csv(test_y_predictions)
```

## 2.4.3 Using Random Forest Regressor

```
In [ ]: rf = RandomForestRegressor(max_depth=30)
         rf.fit(trainingData_X, trainingData_Y)
         rf_predict = rf.predict(test_X)
         write_to_submission_csv(rf_predict)

In [65]: import lightgbm as lgbm

         params = {

             'max_depth': 10
         }

         train_set = lgbm.Dataset(trainingData_X, trainingData_Y, silent=True)
         model = lgbm.train(params, train_set = train_set, num_boost_round=300)
         pred_test_y = model.predict(test_X, num_iteration = model.best_iteration)
         submission = pd.DataFrame(
             {'key': testDF.key, 'fare_amount': pred_test_y},
             columns = ['key', 'fare_amount'])
         submission.to_csv('submission.csv', index = False)
```

## 2.5 Error Metrics

### Error metrics for Linear Regression

```
In [ ]: #Calculating RMSE
         test_X = create_input_features(validateData)
         # Predict fare_amount on the test set using our model (w) trained on the training set.
         test_y_predictions = np.matmul(test_X, weight).round(decimals = 2)
         rmse = sqrt(mean_squared_error(validateData.fare_amount, test_y_predictions))
         rmse
```

## Error Metrics for RandomRegressor LGBM

```
In [ ]: #Calculating RMSE
        test_X = create_input_features(validateData)
        # Predict fare_amount on the test set using our model (w) trained on the training set
        pred_test_y = model.predict(test_X, num_iteration = model.best_iteration)
        rmse = sqrt(mean_squared_error(validateData.fare_amount, pred_test_y))
        rmse
```

## Error metrics for RandomRegressor

```
In [ ]: #Calculating RMSE
        test_X = create_input_features(validateData)
        # Predict fare_amount on the test set using our model (w) trained on the training set.
        rf_predict = rf.predict(test_X)
        rmse = sqrt(mean_squared_error(validateData.fare_amount, rf_predict))
        rmse
```